Running a CEPH-Cluster on a containerized infrastructure

Use case: distributed mySQL-database

Julius Neudecker
Bachelor of Science
julius.neudecker@haw-hamburg.de

January 2020

Contents

Setting up and operate a storage cluster with high availability is a complex task. Modern paradigmas like containerization and orchestration are a way of abstracting away some complexity. However, running a cluster in a stateless and ephemeral containerized environment poses some problems. In the following paper these problems are identified and scrutinized. The use case will be a mySQL database, which will be stored on a CEPH cluster comprised of docker based daemons.